

# resnet\_explore

April 22, 2024

## 1 Imported Modules

```
[ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from torchvision.models import resnet50, ResNet50_Weights
import torchvision.transforms as T

import os
import math as m
from PIL import Image as I
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix as cm

from ray import train, tune
from ray.tune.schedulers import ASHAScheduler
from ray.train import Checkpoint
from ray.tune.search.optuna import OptunaSearch
```

## 2 Model Initialization, Layer Replacement, Device Location

```
[ ]: x = resnet50(weights=ResNet50_Weights.DEFAULT)

for param in x.parameters():
    param.requires_grad = False

x.fc = nn.Linear(2048, 12)
nn.init.kaiming_uniform_(x.fc.weight, nonlinearity='relu')
x.fc.requires_grad = True

if torch.cuda.is_available():
    device = "cuda"
```

```

elif getattr(torch, 'has_mps', False):
    device = "mps"
else:
    device = "cpu"

x = x.to(device)

```

```

/var/folders/yj/n5srfixb95115x6hcn4jcx3000000gn/T/ipykernel_53838/4026304798.py:1
2: UserWarning: 'has_mps' is deprecated, please use
'torch.backends.mps.is_built()'
    elif getattr(torch, 'has_mps', False):

```

### 3 Unnormalization Class

```

[ ]: class Unnormalize(object):
    def __init__(self, mean=None, std=None):
        self.mean = mean
        self.std = std

    def __call__(self, tensor):
        for t, m, s in zip(tensor, self.mean, self.std):
            t.mul_(s).add_(m)
        return tensor

```

### 4 Data Loading Function

```

[ ]: data_labels = {}
    unnorm = Unnormalize()

    def load_images():
        #Bunch of arrays for storing and moving data
        aux_img = []
        aux_labels = []
        train_imgs = []
        train_labels = []
        test_imgs = []
        test_labels = []
        val_imgs = []
        val_labels = []
        idx = 0

        crop = T.CenterCrop((128, 128))

        #Cycle through various folders for images
        for folder in range(len(sorted(os.listdir('./Data')))):
            if sorted(os.listdir('./Data'))[folder] == '.DS_Store':

```

```

        continue
    else:
        aux_img = os.listdir('./Data/'+sorted(os.listdir('./Data'))[folder])

        for i in range(len(aux_img)):
            #Remove .DS_Store file from data
            if aux_img[i] == '.DS_Store':
                aux_img.pop(i)
                break

            #Resize Image and Convert Channels
            for i in range(len(aux_img)):
                try:
                    with I.open('./Data/'+sorted(os.listdir('./
↵Data'))[folder]+'/' +aux_img[i]) as img:
                        img = crop(img)

                        if img.mode != 'RGB':
                            img = img.convert('RGB')

                        img = np.array(img)
                        aux_img[i] = img
                except Exception as e:
                    print(f"{i}")
            aux_labels = []

            #Generate Labels
            for i in range(len(aux_img)):
                aux_labels.append(idx)
            idx += 1

            #Divide Data and Labels into Training, Testing, and Validation Sets
            train_imgs += aux_img[:int(m.ceil(len(aux_img)*0.6))]
            aux_img = aux_img[int(m.ceil(len(aux_img)*0.6)):]
            val_imgs += aux_img[:int(m.ceil(len(aux_img)*0.75))]
            aux_img = aux_img[int(m.ceil(len(aux_img)*0.75)):]
            test_imgs += aux_img

            train_labels += aux_labels[:int(m.ceil(len(aux_labels)*0.6))]
            aux_labels = aux_labels[int(m.ceil(len(aux_labels)*0.6)):]
            val_labels += aux_labels[:int(m.ceil(len(aux_labels)*0.75))]
            aux_labels = aux_labels[int(m.ceil(len(aux_labels)*0.75)):]
            test_labels += aux_labels
            data_labels[folder-1] = sorted(os.listdir('./Data'))[folder]

        #Data Normalization RGB -> [0, 1]
        aux_img = []

```

```

aux_img = train_imgs + test_imgs + val_imgs
mu = np.mean(aux_img, axis=(0, 1, 2))
sigma = np.std(aux_img, axis=(0, 1, 2))
unnorm.mean = mu
unnorm.std = sigma

normalize = T.Compose([
    T.Normalize(mu, sigma)
])

#Conversion to tensors and recast data
train_imgs, test_imgs, val_imgs = torch.tensor(train_imgs).to(torch.
↪float32), torch.tensor(test_imgs).to(torch.float32), torch.tensor(val_imgs).
↪to(torch.float32)
train_labels, test_labels, val_labels = torch.tensor(train_labels), torch.
↪tensor(test_labels), torch.tensor(val_labels)

#Shuffle Images Sets
train_imgs, train_labels = shuffle_imgs(train_imgs, train_labels)
test_imgs, test_labels = shuffle_imgs(test_imgs, test_labels)
val_imgs, val_labels = shuffle_imgs(val_imgs, val_labels)

#Normalize and Reorder Data for Processing
train_imgs = normalize(train_imgs.permute(0, 3, 1, 2))
test_imgs = normalize(test_imgs.permute(0, 3, 1, 2))
val_imgs = normalize(val_imgs.permute(0, 3, 1, 2))

train_set, test_set, val_set = TensorDataset(train_imgs, train_labels),
↪TensorDataset(test_imgs, test_labels), TensorDataset(val_imgs, val_labels)

trainloader, testloader, valloader = DataLoader(dataset=train_set,
↪shuffle=True, batch_size=128), DataLoader(dataset=test_set, shuffle=True,
↪batch_size=64), DataLoader(dataset=val_set, shuffle=True, batch_size=64)

return trainloader, testloader, valloader

#Shuffle Images according to Random Noise Generation
def shuffle_imgs(imgs, labels):
    B = imgs.size(0)
    shuffle = torch.randperm(B)
    return imgs[shuffle], labels[shuffle]

```

```

[ ]: trainloader, testloader, valloader = load_images()
print(data_labels)

```

/var/folders/yj/n5srfixb95115x6hcn4jcx3000000gn/T/ipykernel\_53838/1063623792.py:7  
8: UserWarning: Creating a tensor from a list of numpy.ndarrays is extremely

slow. Please consider converting the list to a single `numpy.ndarray` with `numpy.array()` before converting to a tensor. (Triggered internally at `/Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/utils/tensor_new.cpp:278.`)

```
train_imgs, test_imgs, val_imgs = torch.tensor(train_imgs).to(torch.float32),
torch.tensor(test_imgs).to(torch.float32),
torch.tensor(val_imgs).to(torch.float32)

{0: 'apples', 1: 'coconuts', 2: 'grapes', 3: 'guavas', 4: 'lemons', 5: 'limes',
6: 'mango', 7: 'oranges', 8: 'peaches', 9: 'starfruit', 10: 'strawberries', 11:
'watermelons'}
```

## 5 Train, Test, and Finetuning Functions

```
[ ]: #Training Method with Adam Optimizer
def trainfunc(model, traindata, valdata, epochs):
    optimizer = optim.Adam(model.parameters(), lr=0.006349288161821683, eps=6.
↪5131448417080975e-09)
    #optimizer = optim.SGD(model.parameters(), lr=0.006349288161821683)
    lr_optim = optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=10)
    loss_list = []
    loss_list_aux = []
    validation_loss = []
    validation_loss_aux = []

    for i in range(epochs):
        tot = 0
        div = 0
        for imgs, labels in traindata:
            imgs = imgs.to(device)
            labels = labels.to(device)
            optimizer.zero_grad()

            logits = model(imgs)
            loss = F.cross_entropy(logits, labels)
            loss_list_aux.append(loss.item())

            tot += (torch.argmax(logits, axis=1) == labels).float().sum().item()
            div += imgs.shape[0]

            loss.backward()
            optimizer.step()
        print('Training Accuracy (epoch ', i, '): ', tot/div*100, '%')
        print('Mean Training Loss: ', torch.mean(torch.tensor(loss_list_aux)))
        loss_list.append(torch.mean(torch.tensor(loss_list_aux)))
        lr_optim.step()

    with torch.no_grad():
```

```

        div = 0
        tot = 0
        for imgs, labels in valdata:
            imgs = imgs.to(device)
            labels = labels.to(device)
            val_logits = model(imgs)
            validation_loss_aux.append(F.cross_entropy(val_logits, labels).
↪item())

            val_logits = nn.Softmax()(val_logits)
            pred = torch.argmax(val_logits, axis=1)
            tot += (pred == labels).float().sum().item()
            div += imgs.shape[0]
            validation_loss.append(torch.mean(torch.
↪tensor(validation_loss_aux)))
            print('Mean Validation Loss: ', torch.mean(torch.
↪tensor(validation_loss_aux)))
            print("Validation Accuracy: ", tot/div*100, "%")

        return model, loss_list, validation_loss

#Finetune Function
def finetune(model, data, valdata, epochs):
    model.fc.requires_grad = False
    model.avgpool.requires_grad = True
    model.layer4.requires_grad = True
    model.layer3.requires_grad = True
    model.layer2.requires_grad = True
    model.layer1.requires_grad = True
    model.maxpool.requires_grad = True
    model.conv1.requires_grad = True

    optimizer = optim.Adam(model.parameters(), lr=0.001)
    lr_sched = optim.lr_scheduler.StepLR(optimizer, step_size=2)
    loss_list = []
    loss_list_aux = []

    for i in range(epochs):
        for imgs, labels in data:
            imgs = imgs.to(device)
            labels = labels.to(device)
            optimizer.zero_grad()

            logits = model(imgs)
            loss = F.cross_entropy(logits, labels)
            loss_list_aux.append(loss.item())

            loss.backward()

```

```

        optimizer.step()

    loss_list.append(torch.mean(torch.tensor(loss_list_aux)))

    div = 0
    tot = 0
    for imgs, labels in valdata:
        imgs = imgs.to(device)
        labels = labels.to(device)
        val_logits = model(imgs)
        val_logits = nn.Softmax()(val_logits)
        pred = torch.argmax(val_logits, axis=1)
        tot += (pred == labels).float().sum().item()
        div += imgs.shape[0]
    print("Validation Accuracy: ", tot/div*100, "%")
    lr_sched.step()
    return model, loss_list

def test(model, data):
    div = 0
    total = 0

    for imgs, labels in data:
        imgs = imgs.to(device)
        labels = labels.to(device)
        logits = model(imgs)
        logits = nn.Softmax()(logits)
        pred = torch.argmax(logits, axis=1)
        div += imgs.shape[0]
        total += (pred == labels).float().sum().item()
    print("Accuracy: ", total/div*100, "%")

```

## 6 Confusion Matrix Function

```

[ ]: def confusion_matrix(model, data):
    total_pred = []
    total_labels = []

    for imgs, labels in data:
        imgs = imgs.to(device)
        labels = labels.to(device)

        logits = model(imgs)
        logits = nn.Softmax()(logits)
        pred = torch.argmax(logits, axis=1)

```

```

total_pred += list(pred.to('cpu'))
total_labels += list(labels.to('cpu'))

con_mat = cm(total_labels, total_pred)
labels = [data_labels[i] for i in range(len(con_mat))]
plt.figure(figsize=(10, 8))
ax = sns.heatmap(con_mat, annot=True, cmap="inferno", xticklabels=labels,
yticklabels=labels)
plt.title(f'ResNet50 Accuracy')
plt.ylabel('True Labels')
plt.xlabel('Predicted Labels')
plt.show()
plt.clf()

```

## 7 Hyperparameter Tuning Functions

```

[ ]: def hfinetunetrain(model, data, optimizer, lr_sched):
    for imgs, labels in data:
        imgs, labels = imgs.to(device), labels.to(device)

        pred = model(imgs)
        optimizer.zero_grad()
        loss = F.cross_entropy(pred, labels)
        loss.backward()
        optimizer.step()
        lr_sched.step()
    if device == "cuda":
        torch.cuda.empty_cache()
    elif device == "mps":
        torch.mps.empty_cache()

def hfinetunetest(model, data):
    model.eval()
    correct = 0
    total = 0

    for imgs, labels in data:
        imgs, labels = imgs.to(device), labels.to(device)

        outputs = model(imgs)

        pred = torch.argmax(outputs, axis=1)

        total += imgs.shape[0]
        correct += (pred == labels).float().sum().item()
    if device == "cuda":

```



```

        torch.cuda.empty_cache()
    elif device == "mps":
        torch.mps.empty_cache()
    return correct/total

def objective(config):
    train_loader, test_loader, val_loader = load_images()

    model = resnet50(weights=ResNet50_Weights.DEFAULT).to(device)
    for param in x.parameters():
        param.requires_grad = False

    model.fc = nn.Sequential(nn.Linear(2048, config["l1"]),
                             nn.ReLU(),
                             nn.Linear(config["l1"], config["l2"]),
                             nn.ReLU(),
                             nn.Linear(config["l2"], 12)
                             ).to(device)

    nn.init.kaiming_uniform_(x.fc.weight, nonlinearity='relu')
    x.fc.requires_grad = True

    optimizer = optim.Adam(model.parameters(), lr=config["lr"],
    ↪eps=config["eps"])
    lr_sched = optim.lr_scheduler.StepLR(optimizer,
    ↪step_size=config["lr_sched_step"])

    for epoch in range(10):
        hfinetunetrain(model, train_loader, optimizer, lr_sched)
        acc = hfinetunetest(model, val_loader)

        checkpoint = None
        if (epoch+1)%5 == 0:
            torch.save(model.state_dict(),
                        "Users/thomassigler/DeepLearningProject/
    ↪TuningCheckpoints/model.pth")
            checkpoint = Checkpoint.from_directory("Users/thomassigler/
    ↪DeepLearningProject/TuningCheckpoints")
            train.report({"mean_accuracy": acc}, checkpoint=checkpoint)

        if device == "cuda":
            torch.cuda.empty_cache()
        elif device == "mps":
            torch.mps.empty_cache()

```

## 8 Model Benchmark with Replaced Layer


```
[ ]: test(x, testloader)

confusion_matrix(x, testloader)
```

/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-packages/torch/nn/modules/module.py:1511: UserWarning: Implicit dimension choice for softmax has been deprecated. Change the call to include dim=X as an argument.

```
return self._call_impl(*args, **kwargs)
```

Accuracy: 8.108108108108109 %



<Figure size 640x480 with 0 Axes>

## 9 Training ResNet50 FC Layer

```
[ ]: x, loss_list, val_loss = trainfunc(x, trainloader, valloader, 200)
      test(x, testloader)

      confusion_matrix(x, testloader)
```

```
Training Accuracy (epoch 0 ): 36.67745415318231 %
Mean Training Loss:  tensor(1.8886)
```

```
/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-
packages/torch/nn/modules/module.py:1511: UserWarning: Implicit dimension choice
for softmax has been deprecated. Change the call to include dim=X as an
argument.
```

```
    return self._call_impl(*args, **kwargs)
```

```
Mean Validation Loss:  tensor(1.2270)
Validation Accuracy:  64.64208242950107 %
Training Accuracy (epoch 1 ): 84.46601941747572 %
Mean Training Loss:  tensor(1.3071)
Mean Validation Loss:  tensor(1.0424)
Validation Accuracy:  74.62039045553145 %
Training Accuracy (epoch 2 ): 91.8015102481122 %
Mean Training Loss:  tensor(0.9962)
Mean Validation Loss:  tensor(0.9315)
Validation Accuracy:  77.87418655097615 %
Training Accuracy (epoch 3 ): 96.87162891046385 %
Mean Training Loss:  tensor(0.8044)
Mean Validation Loss:  tensor(0.8995)
Validation Accuracy:  79.39262472885032 %
Training Accuracy (epoch 4 ): 98.70550161812298 %
Mean Training Loss:  tensor(0.6768)
Mean Validation Loss:  tensor(0.8583)
Validation Accuracy:  80.4772234273319 %
Training Accuracy (epoch 5 ): 99.67637540453075 %
Mean Training Loss:  tensor(0.5845)
Mean Validation Loss:  tensor(0.8309)
Validation Accuracy:  78.52494577006507 %
Training Accuracy (epoch 6 ): 100.0 %
Mean Training Loss:  tensor(0.5173)
Mean Validation Loss:  tensor(0.8053)
Validation Accuracy:  79.82646420824295 %
Training Accuracy (epoch 7 ): 99.78425026968716 %
Mean Training Loss:  tensor(0.4655)
Mean Validation Loss:  tensor(0.7808)
Validation Accuracy:  81.77874186550976 %
Training Accuracy (epoch 8 ): 100.0 %
Mean Training Loss:  tensor(0.4249)
Mean Validation Loss:  tensor(0.7626)
```

Validation Accuracy: 81.34490238611714 %  
Training Accuracy (epoch 9 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.3923)  
Mean Validation Loss: tensor(0.7467)  
Validation Accuracy: 80.26030368763557 %  
Training Accuracy (epoch 10 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.3654)  
Mean Validation Loss: tensor(0.7372)  
Validation Accuracy: 79.39262472885032 %  
Training Accuracy (epoch 11 ): 99.46062567421791 %  
Mean Training Loss: tensor(0.3435)  
Mean Validation Loss: tensor(0.7248)  
Validation Accuracy: 81.12798264642083 %  
Training Accuracy (epoch 12 ): 99.67637540453075 %  
Mean Training Loss: tensor(0.3246)  
Mean Validation Loss: tensor(0.7135)  
Validation Accuracy: 80.26030368763557 %  
Training Accuracy (epoch 13 ): 100.0 %  
Mean Training Loss: tensor(0.3082)  
Mean Validation Loss: tensor(0.7070)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 14 ): 100.0 %  
Mean Training Loss: tensor(0.2936)  
Mean Validation Loss: tensor(0.7053)  
Validation Accuracy: 81.56182212581345 %  
Training Accuracy (epoch 15 ): 99.56850053937433 %  
Mean Training Loss: tensor(0.2820)  
Mean Validation Loss: tensor(0.7036)  
Validation Accuracy: 81.12798264642083 %  
Training Accuracy (epoch 16 ): 99.56850053937433 %  
Mean Training Loss: tensor(0.2707)  
Mean Validation Loss: tensor(0.6995)  
Validation Accuracy: 80.26030368763557 %  
Training Accuracy (epoch 17 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.2595)  
Mean Validation Loss: tensor(0.6935)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 18 ): 100.0 %  
Mean Training Loss: tensor(0.2485)  
Mean Validation Loss: tensor(0.6903)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 19 ): 100.0 %  
Mean Training Loss: tensor(0.2384)  
Mean Validation Loss: tensor(0.6895)  
Validation Accuracy: 81.12798264642083 %  
Training Accuracy (epoch 20 ): 100.0 %  
Mean Training Loss: tensor(0.2289)  
Mean Validation Loss: tensor(0.6876)

Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 21 ): 100.0 %  
Mean Training Loss: tensor(0.2202)  
Mean Validation Loss: tensor(0.6836)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 22 ): 100.0 %  
Mean Training Loss: tensor(0.2121)  
Mean Validation Loss: tensor(0.6791)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 23 ): 100.0 %  
Mean Training Loss: tensor(0.2045)  
Mean Validation Loss: tensor(0.6757)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 24 ): 100.0 %  
Mean Training Loss: tensor(0.1973)  
Mean Validation Loss: tensor(0.6728)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 25 ): 100.0 %  
Mean Training Loss: tensor(0.1908)  
Mean Validation Loss: tensor(0.6683)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 26 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.1850)  
Mean Validation Loss: tensor(0.6650)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 27 ): 100.0 %  
Mean Training Loss: tensor(0.1793)  
Mean Validation Loss: tensor(0.6602)  
Validation Accuracy: 84.3817787418655 %  
Training Accuracy (epoch 28 ): 100.0 %  
Mean Training Loss: tensor(0.1738)  
Mean Validation Loss: tensor(0.6602)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 29 ): 100.0 %  
Mean Training Loss: tensor(0.1689)  
Mean Validation Loss: tensor(0.6595)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 30 ): 100.0 %  
Mean Training Loss: tensor(0.1643)  
Mean Validation Loss: tensor(0.6604)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 31 ): 100.0 %  
Mean Training Loss: tensor(0.1599)  
Mean Validation Loss: tensor(0.6596)  
Validation Accuracy: 84.3817787418655 %  
Training Accuracy (epoch 32 ): 100.0 %  
Mean Training Loss: tensor(0.1557)  
Mean Validation Loss: tensor(0.6573)

Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 33 ): 100.0 %  
Mean Training Loss: tensor(0.1517)  
Mean Validation Loss: tensor(0.6552)  
Validation Accuracy: 84.1648590021692 %  
Training Accuracy (epoch 34 ): 100.0 %  
Mean Training Loss: tensor(0.1480)  
Mean Validation Loss: tensor(0.6536)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 35 ): 100.0 %  
Mean Training Loss: tensor(0.1445)  
Mean Validation Loss: tensor(0.6506)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 36 ): 100.0 %  
Mean Training Loss: tensor(0.1412)  
Mean Validation Loss: tensor(0.6496)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 37 ): 99.78425026968716 %  
Mean Training Loss: tensor(0.1385)  
Mean Validation Loss: tensor(0.6486)  
Validation Accuracy: 84.81561822125813 %  
Training Accuracy (epoch 38 ): 100.0 %  
Mean Training Loss: tensor(0.1355)  
Mean Validation Loss: tensor(0.6483)  
Validation Accuracy: 81.56182212581345 %  
Training Accuracy (epoch 39 ): 100.0 %  
Mean Training Loss: tensor(0.1327)  
Mean Validation Loss: tensor(0.6478)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 40 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.1300)  
Mean Validation Loss: tensor(0.6466)  
Validation Accuracy: 81.34490238611714 %  
Training Accuracy (epoch 41 ): 100.0 %  
Mean Training Loss: tensor(0.1275)  
Mean Validation Loss: tensor(0.6443)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 42 ): 100.0 %  
Mean Training Loss: tensor(0.1249)  
Mean Validation Loss: tensor(0.6418)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 43 ): 100.0 %  
Mean Training Loss: tensor(0.1224)  
Mean Validation Loss: tensor(0.6395)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 44 ): 100.0 %  
Mean Training Loss: tensor(0.1200)  
Mean Validation Loss: tensor(0.6374)

Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 45 ): 100.0 %  
Mean Training Loss: tensor(0.1177)  
Mean Validation Loss: tensor(0.6370)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 46 ): 100.0 %  
Mean Training Loss: tensor(0.1155)  
Mean Validation Loss: tensor(0.6394)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 47 ): 100.0 %  
Mean Training Loss: tensor(0.1134)  
Mean Validation Loss: tensor(0.6373)  
Validation Accuracy: 83.73101952277658 %  
Training Accuracy (epoch 48 ): 100.0 %  
Mean Training Loss: tensor(0.1112)  
Mean Validation Loss: tensor(0.6381)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 49 ): 100.0 %  
Mean Training Loss: tensor(0.1092)  
Mean Validation Loss: tensor(0.6368)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 50 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.1075)  
Mean Validation Loss: tensor(0.6358)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 51 ): 100.0 %  
Mean Training Loss: tensor(0.1056)  
Mean Validation Loss: tensor(0.6356)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 52 ): 100.0 %  
Mean Training Loss: tensor(0.1039)  
Mean Validation Loss: tensor(0.6348)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 53 ): 100.0 %  
Mean Training Loss: tensor(0.1022)  
Mean Validation Loss: tensor(0.6385)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 54 ): 100.0 %  
Mean Training Loss: tensor(0.1005)  
Mean Validation Loss: tensor(0.6371)  
Validation Accuracy: 81.34490238611714 %  
Training Accuracy (epoch 55 ): 100.0 %  
Mean Training Loss: tensor(0.0989)  
Mean Validation Loss: tensor(0.6389)  
Validation Accuracy: 81.56182212581345 %  
Training Accuracy (epoch 56 ): 100.0 %  
Mean Training Loss: tensor(0.0974)  
Mean Validation Loss: tensor(0.6390)

Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 57 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.0960)  
Mean Validation Loss: tensor(0.6382)  
Validation Accuracy: 80.91106290672451 %  
Training Accuracy (epoch 58 ): 100.0 %  
Mean Training Loss: tensor(0.0946)  
Mean Validation Loss: tensor(0.6374)  
Validation Accuracy: 83.73101952277658 %  
Training Accuracy (epoch 59 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.0932)  
Mean Validation Loss: tensor(0.6372)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 60 ): 100.0 %  
Mean Training Loss: tensor(0.0919)  
Mean Validation Loss: tensor(0.6378)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 61 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.0906)  
Mean Validation Loss: tensor(0.6363)  
Validation Accuracy: 84.81561822125813 %  
Training Accuracy (epoch 62 ): 100.0 %  
Mean Training Loss: tensor(0.0894)  
Mean Validation Loss: tensor(0.6356)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 63 ): 100.0 %  
Mean Training Loss: tensor(0.0881)  
Mean Validation Loss: tensor(0.6351)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 64 ): 100.0 %  
Mean Training Loss: tensor(0.0869)  
Mean Validation Loss: tensor(0.6356)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 65 ): 100.0 %  
Mean Training Loss: tensor(0.0857)  
Mean Validation Loss: tensor(0.6346)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 66 ): 100.0 %  
Mean Training Loss: tensor(0.0847)  
Mean Validation Loss: tensor(0.6347)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 67 ): 100.0 %  
Mean Training Loss: tensor(0.0836)  
Mean Validation Loss: tensor(0.6341)  
Validation Accuracy: 83.73101952277658 %  
Training Accuracy (epoch 68 ): 100.0 %  
Mean Training Loss: tensor(0.0825)  
Mean Validation Loss: tensor(0.6347)



Validation Accuracy: 83.73101952277658 %  
Training Accuracy (epoch 69 ): 100.0 %  
Mean Training Loss: tensor(0.0815)  
Mean Validation Loss: tensor(0.6339)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 70 ): 100.0 %  
Mean Training Loss: tensor(0.0804)  
Mean Validation Loss: tensor(0.6341)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 71 ): 100.0 %  
Mean Training Loss: tensor(0.0794)  
Mean Validation Loss: tensor(0.6336)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 72 ): 100.0 %  
Mean Training Loss: tensor(0.0785)  
Mean Validation Loss: tensor(0.6335)  
Validation Accuracy: 81.12798264642083 %  
Training Accuracy (epoch 73 ): 100.0 %  
Mean Training Loss: tensor(0.0775)  
Mean Validation Loss: tensor(0.6329)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 74 ): 100.0 %  
Mean Training Loss: tensor(0.0766)  
Mean Validation Loss: tensor(0.6323)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 75 ): 100.0 %  
Mean Training Loss: tensor(0.0756)  
Mean Validation Loss: tensor(0.6320)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 76 ): 100.0 %  
Mean Training Loss: tensor(0.0748)  
Mean Validation Loss: tensor(0.6319)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 77 ): 100.0 %  
Mean Training Loss: tensor(0.0739)  
Mean Validation Loss: tensor(0.6314)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 78 ): 100.0 %  
Mean Training Loss: tensor(0.0731)  
Mean Validation Loss: tensor(0.6314)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 79 ): 100.0 %  
Mean Training Loss: tensor(0.0723)  
Mean Validation Loss: tensor(0.6312)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 80 ): 100.0 %  
Mean Training Loss: tensor(0.0715)  
Mean Validation Loss: tensor(0.6314)

Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 81 ): 100.0 %  
Mean Training Loss: tensor(0.0707)  
Mean Validation Loss: tensor(0.6304)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 82 ): 100.0 %  
Mean Training Loss: tensor(0.0699)  
Mean Validation Loss: tensor(0.6302)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 83 ): 100.0 %  
Mean Training Loss: tensor(0.0692)  
Mean Validation Loss: tensor(0.6306)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 84 ): 100.0 %  
Mean Training Loss: tensor(0.0685)  
Mean Validation Loss: tensor(0.6305)  
Validation Accuracy: 84.1648590021692 %  
Training Accuracy (epoch 85 ): 100.0 %  
Mean Training Loss: tensor(0.0678)  
Mean Validation Loss: tensor(0.6305)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 86 ): 100.0 %  
Mean Training Loss: tensor(0.0670)  
Mean Validation Loss: tensor(0.6309)  
Validation Accuracy: 83.73101952277658 %  
Training Accuracy (epoch 87 ): 100.0 %  
Mean Training Loss: tensor(0.0664)  
Mean Validation Loss: tensor(0.6304)  
Validation Accuracy: 83.94793926247289 %  
Training Accuracy (epoch 88 ): 100.0 %  
Mean Training Loss: tensor(0.0657)  
Mean Validation Loss: tensor(0.6298)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 89 ): 100.0 %  
Mean Training Loss: tensor(0.0650)  
Mean Validation Loss: tensor(0.6288)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 90 ): 100.0 %  
Mean Training Loss: tensor(0.0644)  
Mean Validation Loss: tensor(0.6290)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 91 ): 100.0 %  
Mean Training Loss: tensor(0.0637)  
Mean Validation Loss: tensor(0.6276)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 92 ): 100.0 %  
Mean Training Loss: tensor(0.0631)  
Mean Validation Loss: tensor(0.6278)

Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 93 ): 100.0 %  
Mean Training Loss: tensor(0.0625)  
Mean Validation Loss: tensor(0.6275)  
Validation Accuracy: 84.1648590021692 %  
Training Accuracy (epoch 94 ): 100.0 %  
Mean Training Loss: tensor(0.0619)  
Mean Validation Loss: tensor(0.6272)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 95 ): 100.0 %  
Mean Training Loss: tensor(0.0613)  
Mean Validation Loss: tensor(0.6260)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 96 ): 99.78425026968716 %  
Mean Training Loss: tensor(0.0608)  
Mean Validation Loss: tensor(0.6254)  
Validation Accuracy: 84.81561822125813 %  
Training Accuracy (epoch 97 ): 100.0 %  
Mean Training Loss: tensor(0.0602)  
Mean Validation Loss: tensor(0.6249)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 98 ): 100.0 %  
Mean Training Loss: tensor(0.0597)  
Mean Validation Loss: tensor(0.6260)  
Validation Accuracy: 81.56182212581345 %  
Training Accuracy (epoch 99 ): 100.0 %  
Mean Training Loss: tensor(0.0592)  
Mean Validation Loss: tensor(0.6261)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 100 ): 100.0 %  
Mean Training Loss: tensor(0.0586)  
Mean Validation Loss: tensor(0.6264)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 101 ): 100.0 %  
Mean Training Loss: tensor(0.0581)  
Mean Validation Loss: tensor(0.6263)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 102 ): 100.0 %  
Mean Training Loss: tensor(0.0576)  
Mean Validation Loss: tensor(0.6261)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 103 ): 100.0 %  
Mean Training Loss: tensor(0.0570)  
Mean Validation Loss: tensor(0.6261)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 104 ): 100.0 %  
Mean Training Loss: tensor(0.0565)  
Mean Validation Loss: tensor(0.6258)

Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 105 ): 100.0 %  
Mean Training Loss: tensor(0.0560)  
Mean Validation Loss: tensor(0.6254)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 106 ): 100.0 %  
Mean Training Loss: tensor(0.0556)  
Mean Validation Loss: tensor(0.6252)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 107 ): 100.0 %  
Mean Training Loss: tensor(0.0551)  
Mean Validation Loss: tensor(0.6270)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 108 ): 100.0 %  
Mean Training Loss: tensor(0.0546)  
Mean Validation Loss: tensor(0.6265)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 109 ): 100.0 %  
Mean Training Loss: tensor(0.0541)  
Mean Validation Loss: tensor(0.6266)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 110 ): 100.0 %  
Mean Training Loss: tensor(0.0537)  
Mean Validation Loss: tensor(0.6264)  
Validation Accuracy: 84.1648590021692 %  
Training Accuracy (epoch 111 ): 100.0 %  
Mean Training Loss: tensor(0.0533)  
Mean Validation Loss: tensor(0.6255)  
Validation Accuracy: 84.81561822125813 %  
Training Accuracy (epoch 112 ): 100.0 %  
Mean Training Loss: tensor(0.0528)  
Mean Validation Loss: tensor(0.6261)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 113 ): 100.0 %  
Mean Training Loss: tensor(0.0524)  
Mean Validation Loss: tensor(0.6259)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 114 ): 100.0 %  
Mean Training Loss: tensor(0.0520)  
Mean Validation Loss: tensor(0.6256)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 115 ): 100.0 %  
Mean Training Loss: tensor(0.0516)  
Mean Validation Loss: tensor(0.6257)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 116 ): 100.0 %  
Mean Training Loss: tensor(0.0512)  
Mean Validation Loss: tensor(0.6253)

Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 117 ): 100.0 %  
Mean Training Loss: tensor(0.0508)  
Mean Validation Loss: tensor(0.6253)  
Validation Accuracy: 82.21258134490239 %  
Training Accuracy (epoch 118 ): 100.0 %  
Mean Training Loss: tensor(0.0504)  
Mean Validation Loss: tensor(0.6251)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 119 ): 99.89212513484358 %  
Mean Training Loss: tensor(0.0501)  
Mean Validation Loss: tensor(0.6254)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 120 ): 100.0 %  
Mean Training Loss: tensor(0.0497)  
Mean Validation Loss: tensor(0.6264)  
Validation Accuracy: 80.26030368763557 %  
Training Accuracy (epoch 121 ): 100.0 %  
Mean Training Loss: tensor(0.0493)  
Mean Validation Loss: tensor(0.6266)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 122 ): 100.0 %  
Mean Training Loss: tensor(0.0490)  
Mean Validation Loss: tensor(0.6266)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 123 ): 100.0 %  
Mean Training Loss: tensor(0.0486)  
Mean Validation Loss: tensor(0.6266)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 124 ): 100.0 %  
Mean Training Loss: tensor(0.0482)  
Mean Validation Loss: tensor(0.6265)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 125 ): 100.0 %  
Mean Training Loss: tensor(0.0479)  
Mean Validation Loss: tensor(0.6265)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 126 ): 100.0 %  
Mean Training Loss: tensor(0.0475)  
Mean Validation Loss: tensor(0.6261)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 127 ): 100.0 %  
Mean Training Loss: tensor(0.0472)  
Mean Validation Loss: tensor(0.6263)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 128 ): 100.0 %  
Mean Training Loss: tensor(0.0469)  
Mean Validation Loss: tensor(0.6258)

Validation Accuracy: 84.3817787418655 %  
Training Accuracy (epoch 129 ): 100.0 %  
Mean Training Loss: tensor(0.0465)  
Mean Validation Loss: tensor(0.6258)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 130 ): 100.0 %  
Mean Training Loss: tensor(0.0462)  
Mean Validation Loss: tensor(0.6255)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 131 ): 100.0 %  
Mean Training Loss: tensor(0.0459)  
Mean Validation Loss: tensor(0.6257)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 132 ): 100.0 %  
Mean Training Loss: tensor(0.0456)  
Mean Validation Loss: tensor(0.6254)  
Validation Accuracy: 84.59869848156181 %  
Training Accuracy (epoch 133 ): 100.0 %  
Mean Training Loss: tensor(0.0453)  
Mean Validation Loss: tensor(0.6254)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 134 ): 100.0 %  
Mean Training Loss: tensor(0.0450)  
Mean Validation Loss: tensor(0.6252)  
Validation Accuracy: 83.73101952277658 %  
Training Accuracy (epoch 135 ): 100.0 %  
Mean Training Loss: tensor(0.0446)  
Mean Validation Loss: tensor(0.6258)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 136 ): 100.0 %  
Mean Training Loss: tensor(0.0443)  
Mean Validation Loss: tensor(0.6257)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 137 ): 100.0 %  
Mean Training Loss: tensor(0.0440)  
Mean Validation Loss: tensor(0.6255)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 138 ): 100.0 %  
Mean Training Loss: tensor(0.0437)  
Mean Validation Loss: tensor(0.6253)  
Validation Accuracy: 84.1648590021692 %  
Training Accuracy (epoch 139 ): 100.0 %  
Mean Training Loss: tensor(0.0435)  
Mean Validation Loss: tensor(0.6252)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 140 ): 100.0 %  
Mean Training Loss: tensor(0.0432)  
Mean Validation Loss: tensor(0.6251)

Validation Accuracy: 84.1648590021692 %  
Training Accuracy (epoch 141 ): 100.0 %  
Mean Training Loss: tensor(0.0429)  
Mean Validation Loss: tensor(0.6247)  
Validation Accuracy: 81.99566160520607 %  
Training Accuracy (epoch 142 ): 100.0 %  
Mean Training Loss: tensor(0.0426)  
Mean Validation Loss: tensor(0.6250)  
Validation Accuracy: 81.77874186550976 %  
Training Accuracy (epoch 143 ): 100.0 %  
Mean Training Loss: tensor(0.0423)  
Mean Validation Loss: tensor(0.6251)  
Validation Accuracy: 80.91106290672451 %  
Training Accuracy (epoch 144 ): 100.0 %  
Mean Training Loss: tensor(0.0421)  
Mean Validation Loss: tensor(0.6257)  
Validation Accuracy: 83.94793926247289 %  
Training Accuracy (epoch 145 ): 100.0 %  
Mean Training Loss: tensor(0.0418)  
Mean Validation Loss: tensor(0.6255)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 146 ): 100.0 %  
Mean Training Loss: tensor(0.0416)  
Mean Validation Loss: tensor(0.6266)  
Validation Accuracy: 81.77874186550976 %  
Training Accuracy (epoch 147 ): 100.0 %  
Mean Training Loss: tensor(0.0413)  
Mean Validation Loss: tensor(0.6264)  
Validation Accuracy: 84.59869848156181 %  
Training Accuracy (epoch 148 ): 100.0 %  
Mean Training Loss: tensor(0.0411)  
Mean Validation Loss: tensor(0.6270)  
Validation Accuracy: 81.56182212581345 %  
Training Accuracy (epoch 149 ): 100.0 %  
Mean Training Loss: tensor(0.0408)  
Mean Validation Loss: tensor(0.6267)  
Validation Accuracy: 83.94793926247289 %  
Training Accuracy (epoch 150 ): 100.0 %  
Mean Training Loss: tensor(0.0406)  
Mean Validation Loss: tensor(0.6266)  
Validation Accuracy: 83.73101952277658 %  
Training Accuracy (epoch 151 ): 100.0 %  
Mean Training Loss: tensor(0.0403)  
Mean Validation Loss: tensor(0.6273)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 152 ): 100.0 %  
Mean Training Loss: tensor(0.0401)  
Mean Validation Loss: tensor(0.6277)


Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 153 ): 100.0 %  
Mean Training Loss: tensor(0.0398)  
Mean Validation Loss: tensor(0.6281)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 154 ): 100.0 %  
Mean Training Loss: tensor(0.0396)  
Mean Validation Loss: tensor(0.6282)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 155 ): 100.0 %  
Mean Training Loss: tensor(0.0394)  
Mean Validation Loss: tensor(0.6280)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 156 ): 100.0 %  
Mean Training Loss: tensor(0.0391)  
Mean Validation Loss: tensor(0.6275)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 157 ): 100.0 %  
Mean Training Loss: tensor(0.0389)  
Mean Validation Loss: tensor(0.6273)  
Validation Accuracy: 81.12798264642083 %  
Training Accuracy (epoch 158 ): 100.0 %  
Mean Training Loss: tensor(0.0387)  
Mean Validation Loss: tensor(0.6273)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 159 ): 100.0 %  
Mean Training Loss: tensor(0.0384)  
Mean Validation Loss: tensor(0.6272)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 160 ): 100.0 %  
Mean Training Loss: tensor(0.0382)  
Mean Validation Loss: tensor(0.6272)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 161 ): 100.0 %  
Mean Training Loss: tensor(0.0380)  
Mean Validation Loss: tensor(0.6274)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 162 ): 100.0 %  
Mean Training Loss: tensor(0.0378)  
Mean Validation Loss: tensor(0.6271)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 163 ): 100.0 %  
Mean Training Loss: tensor(0.0376)  
Mean Validation Loss: tensor(0.6271)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 164 ): 100.0 %  
Mean Training Loss: tensor(0.0373)  
Mean Validation Loss: tensor(0.6268)



Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 165 ): 100.0 %  
Mean Training Loss: tensor(0.0371)  
Mean Validation Loss: tensor(0.6268)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 166 ): 100.0 %  
Mean Training Loss: tensor(0.0369)  
Mean Validation Loss: tensor(0.6267)  
Validation Accuracy: 83.73101952277658 %  
Training Accuracy (epoch 167 ): 100.0 %  
Mean Training Loss: tensor(0.0367)  
Mean Validation Loss: tensor(0.6264)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 168 ): 100.0 %  
Mean Training Loss: tensor(0.0365)  
Mean Validation Loss: tensor(0.6269)  
Validation Accuracy: 84.3817787418655 %  
Training Accuracy (epoch 169 ): 100.0 %  
Mean Training Loss: tensor(0.0363)  
Mean Validation Loss: tensor(0.6270)  
Validation Accuracy: 83.08026030368764 %  
Training Accuracy (epoch 170 ): 100.0 %  
Mean Training Loss: tensor(0.0361)  
Mean Validation Loss: tensor(0.6269)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 171 ): 100.0 %  
Mean Training Loss: tensor(0.0359)  
Mean Validation Loss: tensor(0.6271)  
Validation Accuracy: 83.94793926247289 %  
Training Accuracy (epoch 172 ): 100.0 %  
Mean Training Loss: tensor(0.0357)  
Mean Validation Loss: tensor(0.6272)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 173 ): 100.0 %  
Mean Training Loss: tensor(0.0356)  
Mean Validation Loss: tensor(0.6275)  
Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 174 ): 100.0 %  
Mean Training Loss: tensor(0.0354)  
Mean Validation Loss: tensor(0.6277)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 175 ): 100.0 %  
Mean Training Loss: tensor(0.0352)  
Mean Validation Loss: tensor(0.6278)  
Validation Accuracy: 83.94793926247289 %  
Training Accuracy (epoch 176 ): 100.0 %  
Mean Training Loss: tensor(0.0350)  
Mean Validation Loss: tensor(0.6281)

Validation Accuracy: 82.646420824295 %  
Training Accuracy (epoch 177 ): 100.0 %  
Mean Training Loss: tensor(0.0348)  
Mean Validation Loss: tensor(0.6278)  
Validation Accuracy: 82.86334056399133 %  
Training Accuracy (epoch 178 ): 100.0 %  
Mean Training Loss: tensor(0.0346)  
Mean Validation Loss: tensor(0.6282)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 179 ): 100.0 %  
Mean Training Loss: tensor(0.0345)  
Mean Validation Loss: tensor(0.6282)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 180 ): 100.0 %  
Mean Training Loss: tensor(0.0343)  
Mean Validation Loss: tensor(0.6277)  
Validation Accuracy: 84.1648590021692 %  
Training Accuracy (epoch 181 ): 100.0 %  
Mean Training Loss: tensor(0.0341)  
Mean Validation Loss: tensor(0.6279)  
Validation Accuracy: 84.81561822125813 %  
Training Accuracy (epoch 182 ): 100.0 %  
Mean Training Loss: tensor(0.0339)  
Mean Validation Loss: tensor(0.6280)  
Validation Accuracy: 83.94793926247289 %  
Training Accuracy (epoch 183 ): 100.0 %  
Mean Training Loss: tensor(0.0337)  
Mean Validation Loss: tensor(0.6288)  
Validation Accuracy: 83.94793926247289 %  
Training Accuracy (epoch 184 ): 100.0 %  
Mean Training Loss: tensor(0.0336)  
Mean Validation Loss: tensor(0.6288)  
Validation Accuracy: 83.29718004338395 %  
Training Accuracy (epoch 185 ): 100.0 %  
Mean Training Loss: tensor(0.0334)  
Mean Validation Loss: tensor(0.6289)  
Validation Accuracy: 82.4295010845987 %  
Training Accuracy (epoch 186 ): 100.0 %  
Mean Training Loss: tensor(0.0332)  
Mean Validation Loss: tensor(0.6285)  
Validation Accuracy: 83.51409978308027 %  
Training Accuracy (epoch 187 ): 100.0 %  
Mean Training Loss: tensor(0.0331)  
Mean Validation Loss: tensor(0.6286)  
Validation Accuracy: 81.77874186550976 %  
Training Accuracy (epoch 188 ): 100.0 %  
Mean Training Loss: tensor(0.0329)  
Mean Validation Loss: tensor(0.6283)

Validation Accuracy: 83.29718004338395 %  
 Training Accuracy (epoch 189 ): 100.0 %  
 Mean Training Loss: tensor(0.0328)  
 Mean Validation Loss: tensor(0.6280)  
 Validation Accuracy: 82.86334056399133 %  
 Training Accuracy (epoch 190 ): 100.0 %  
 Mean Training Loss: tensor(0.0326)  
 Mean Validation Loss: tensor(0.6282)  
 Validation Accuracy: 83.51409978308027 %  
 Training Accuracy (epoch 191 ): 100.0 %  
 Mean Training Loss: tensor(0.0324)  
 Mean Validation Loss: tensor(0.6283)  
 Validation Accuracy: 83.08026030368764 %  
 Training Accuracy (epoch 192 ): 100.0 %  
 Mean Training Loss: tensor(0.0323)  
 Mean Validation Loss: tensor(0.6286)  
 Validation Accuracy: 83.51409978308027 %  
 Training Accuracy (epoch 193 ): 100.0 %  
 Mean Training Loss: tensor(0.0321)  
 Mean Validation Loss: tensor(0.6283)  
 Validation Accuracy: 83.94793926247289 %  
 Training Accuracy (epoch 194 ): 100.0 %  
 Mean Training Loss: tensor(0.0320)  
 Mean Validation Loss: tensor(0.6285)  
 Validation Accuracy: 83.94793926247289 %  
 Training Accuracy (epoch 195 ): 100.0 %  
 Mean Training Loss: tensor(0.0318)  
 Mean Validation Loss: tensor(0.6284)  
 Validation Accuracy: 83.73101952277658 %  
 Training Accuracy (epoch 196 ): 100.0 %  
 Mean Training Loss: tensor(0.0317)  
 Mean Validation Loss: tensor(0.6282)  
 Validation Accuracy: 82.21258134490239 %  
 Training Accuracy (epoch 197 ): 99.89212513484358 %  
 Mean Training Loss: tensor(0.0315)  
 Mean Validation Loss: tensor(0.6291)  
 Validation Accuracy: 85.24945770065075 %  
 Training Accuracy (epoch 198 ): 100.0 %  
 Mean Training Loss: tensor(0.0314)  
 Mean Validation Loss: tensor(0.6287)  
 Validation Accuracy: 83.29718004338395 %  
 Training Accuracy (epoch 199 ): 100.0 %  
 Mean Training Loss: tensor(0.0313)  
 Mean Validation Loss: tensor(0.6299)  
 Validation Accuracy: 84.81561822125813 %  
 Accuracy: 81.08108108108108 %



<Figure size 640x480 with 0 Axes>

## 10 Finetuning Pre-trained Layers

```
[ ]: x, loss_list_val = finetune(x, valloader, valloader, 200)
test(x, testloader)
```

/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-packages/torch/nn/modules/module.py:1511: UserWarning: Implicit dimension choice for softmax has been deprecated. Change the call to include dim=X as an argument.

```
return self._call_impl(*args, **kwargs)
```

Validation Accuracy: 85.46637744034707 %

Validation Accuracy: 85.24945770065075 %

Validation Accuracy: 86.55097613882863 %

Validation Accuracy: 84.59869848156181 %  
Validation Accuracy: 85.24945770065075 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 88.06941431670282 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 88.06941431670282 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 83.94793926247289 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 85.03253796095444 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 88.06941431670282 %  
Validation Accuracy: 85.24945770065075 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 87.85249457700651 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 85.46637744034707 %

Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 85.24945770065075 %  
Validation Accuracy: 84.1648590021692 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 84.81561822125813 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 88.28633405639913 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 85.24945770065075 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 84.59869848156181 %  
Validation Accuracy: 87.63557483731019 %

Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 85.24945770065075 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 87.85249457700651 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 87.63557483731019 %  
Validation Accuracy: 84.59869848156181 %  
Validation Accuracy: 85.03253796095444 %  
Validation Accuracy: 85.03253796095444 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 84.1648590021692 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 85.24945770065075 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 87.85249457700651 %  
Validation Accuracy: 87.85249457700651 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 87.63557483731019 %  
Validation Accuracy: 87.63557483731019 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 85.03253796095444 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 86.11713665943601 %

Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 88.28633405639913 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 87.85249457700651 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 87.85249457700651 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 85.68329718004338 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 85.03253796095444 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 84.3817787418655 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 87.41865509761388 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 86.11713665943601 %  
Validation Accuracy: 85.90021691973969 %  
Validation Accuracy: 88.93709327548807 %  
Validation Accuracy: 85.03253796095444 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 86.98481561822126 %  
Validation Accuracy: 87.63557483731019 %  
Validation Accuracy: 86.33405639913232 %  
Validation Accuracy: 84.81561822125813 %  
Validation Accuracy: 86.76789587852495 %  
Validation Accuracy: 85.46637744034707 %  
Validation Accuracy: 86.55097613882863 %  
Validation Accuracy: 87.20173535791757 %  
Validation Accuracy: 85.24945770065075 %  
Validation Accuracy: 87.85249457700651 %



```

Validation Accuracy: 85.68329718004338 %
Validation Accuracy: 86.55097613882863 %
Validation Accuracy: 86.11713665943601 %
Validation Accuracy: 86.33405639913232 %
Validation Accuracy: 86.76789587852495 %
Accuracy: 85.8108108108108 %

```

## 11 Hyperparameter Tuning

Hyper-parameter tuning requires that all file path names in this script be their absolute path names. Otherwise, this portion of the program will crash. This is due to how Raytune runs its testing environment, which is in a different file from this script.

This portion of the script can be slow if CUDA is not used. There is 0 support for Apple Metal by Ray, so the cpu will be used for testing hyperparameters even if items are loaded to MPS.

Ray was built for Linux and Linux-based systems. Running this portion of the script on a Windows device can be difficult and has been a known issue since late 2023. This is due to a fix the library implemented for a Pyarrow 7+ bug.

```

[ ]: search_space = {"lr": tune.loguniform(1e-4, 1e-1), "eps": tune.
    ↳loguniform(1e-10, 1e-6), "l1": tune.randint(50, 2049), "l2": tune.
    ↳randint(12, 101), "lr_sched_step": tune.randint(1, 21)}
tuner = tune.Tuner(objective, tune_config=tune.
    ↳TuneConfig(metric="mean_accuracy", mode="max", search_alg=OptunaSearch(),
    ↳num_samples=2), param_space=search_space)

results = tuner.fit()
print("Best Config is: ", results.get_best_result().config)

```

<IPython.core.display.HTML object>

(raylet) Warning: The actor ImplicitFunc is very large (90 MiB). Check that its definition is not implicitly capturing a large array or other object in scope. Tip: use ray.put() to put large objects in the Ray object store.

(objective pid=50259) /var/folders/yj/n5srfixb951l5x6hcn4jcx3000000gn/T/ipykernel\_51918/3592750841.py:78: UserWarning: Creating a tensor from a list of numpy.ndarrays is extremely slow. Please consider converting the list to a single numpy.ndarray with numpy.array() before converting to a tensor.

(Triggered internally at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/utils/tensor\_new.cpp:278.)

(objective pid=50277) /var/folders/yj/n5srfixb951l5x6hcn4jcx3000000gn/T/ipykernel\_51918/3592750841.py:78: UserWarning: Creating a tensor from a list of numpy.ndarrays is extremely slow. Please consider converting the list to a single numpy.ndarray with numpy.array() before converting to a tensor.

(Triggered internally at /Users/runner/work/pytorch/pytorch/pytorch/torch/csrc/utils/tensor\_new.cpp:278.)

2024-04-22 18:10:08,284 WARNING util.py:202 -- The `on\_step\_begin` operation took 0.550 s, which may be a performance bottleneck.

```

(objective pid=50259) Checkpoint successfully created at:
Checkpoint(filesystem=local, path=/Users/thomassigler/ray_results/objective_2024-04-22_17-59-50/objective_b9f0a4dd_1_eps=0.0000,l1=1881,l2=42,lr=0.0002,lr_sched_step=6_2024-04-22_17-59-55/checkpoint_000000)
(objective pid=50277) Checkpoint successfully created at:
Checkpoint(filesystem=local, path=/Users/thomassigler/ray_results/objective_2024-04-22_17-59-50/objective_f4523408_2_eps=0.0000,l1=719,l2=72,lr=0.0125,lr_sched_step=6_2024-04-22_18-00-11/checkpoint_000000)
(objective pid=50259) Checkpoint successfully created at:
Checkpoint(filesystem=local, path=/Users/thomassigler/ray_results/objective_2024-04-22_17-59-50/objective_b9f0a4dd_1_eps=0.0000,l1=1881,l2=42,lr=0.0002,lr_sched_step=6_2024-04-22_17-59-55/checkpoint_000001)
(objective pid=50277) Checkpoint successfully created at:
Checkpoint(filesystem=local, path=/Users/thomassigler/ray_results/objective_2024-04-22_17-59-50/objective_f4523408_2_eps=0.0000,l1=719,l2=72,lr=0.0125,lr_sched_step=6_2024-04-22_18-00-11/checkpoint_000001)
2024-04-22 18:27:02,663 INFO tune.py:1016 -- Wrote the latest version of all
result files and experiment state to
'/Users/thomassigler/ray_results/objective_2024-04-22_17-59-50' in 0.0138s.
2024-04-22 18:27:02,697 INFO tune.py:1048 -- Total run time: 1632.33 seconds
(1630.50 seconds for the tuning loop).

Best Config is: {'lr': 0.00023301901230003848, 'eps': 2.887185712485118e-09,
'11': 1881, '12': 42, 'lr_sched_step': 6}

```

## 12 Plotting Graphs

```

[ ]: confusion_matrix(x, testloader)

plt.plot(np.linspace(1, len(loss_list), len(loss_list)), loss_list,
         label='Training Loss')
plt.plot(np.linspace(1, len(val_loss), len(loss_list)), loss_list,
         label='Validation Loss')
plt.title('Training and Validation Loss over Epochs No Fine Tuning')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.grid(True)
plt.legend()
plt.show()
plt.clf()


plt.plot(np.linspace(1, len(loss_list_val), len(loss_list_val)), loss_list_val)
plt.title('Loss over Epochs Finetuning')


```

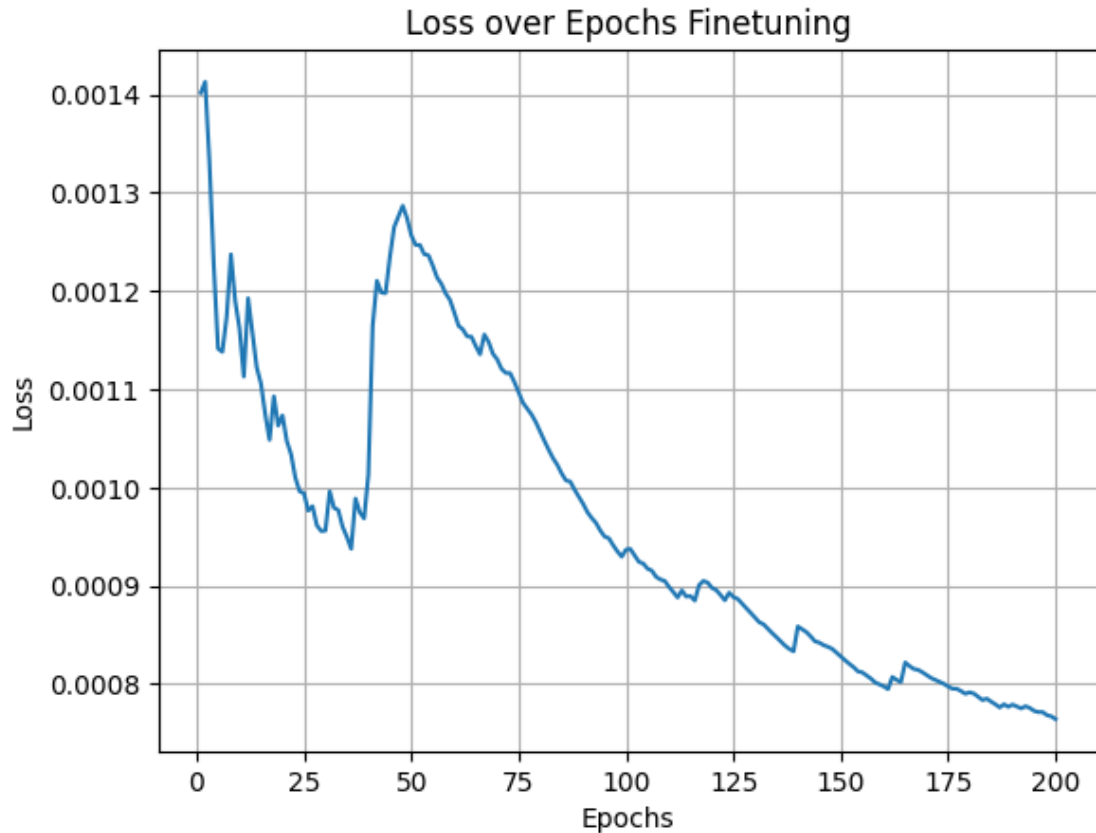
```
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.grid(True)
plt.show()
plt.clf()
```

/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-packages/torch/nn/modules/module.py:1511: UserWarning: Implicit dimension choice for softmax has been deprecated. Change the call to include dim=X as an argument.

```
return self._call_impl(*args, **kwargs)
```







## 13 Optimizer Testing and Visualization

```
[ ]: def optimizer_testing(data, model, optimizer, scheduler, val_data, epochs=100):
    model = model.to("mps")
    acc_list = []
    for epoch in range(epochs):
        count = 0
        total = 0
        for imgs, labels in data:
            imgs, labels = imgs.to(device), labels.to(device)

            logits = model(imgs)
            loss = F.cross_entropy(logits, labels)

            loss.backward()
            optimizer.step()
            scheduler.step()

        for imgs, labels in val_data:
```

```

        imgs, labels = imgs.to(device), labels.to(device)

        logits = model(imgs)
        loss = F.cross_entropy(logits, labels)

        loss.backward()
        optimizer.step()

        pred = torch.argmax(logits, axis=1)
        count += (pred == labels).float().sum().item()
        total += imgs.shape[0]

    acc_list.append(count/total*100)
    return acc_list

```

```

[ ]: aux_model = resnet50(weights=ResNet50_Weights.DEFAULT)

for param in aux_model.parameters():
    param.requires_grad = False

aux_model.fc = nn.Linear(2048, 12)
nn.init.kaiming_uniform_(aux_model.fc.weight, nonlinearity='relu')
aux_model.fc.requires_grad = True
optimizer = torch.optim.Adam(aux_model.parameters(), lr=0.01)
Adam_list_1 = optimizer_testing(trainloader, aux_model, optimizer, optim.
    ↳lr_scheduler.StepLR(optimizer, step_size=1), valloader)
print("Done - 1")
aux_model = resnet50(weights=ResNet50_Weights.DEFAULT)

for param in aux_model.parameters():
    param.requires_grad = False

aux_model.fc = nn.Linear(2048, 12)
nn.init.kaiming_uniform_(aux_model.fc.weight, nonlinearity='relu')
aux_model.fc.requires_grad = True
optimizer = torch.optim.Adam(aux_model.parameters(), lr=0.01)
Adam_list_2 = optimizer_testing(trainloader, aux_model, optimizer, optim.
    ↳lr_scheduler.StepLR(optimizer, step_size=2), valloader)
print("Done - 2")
aux_model = resnet50(weights=ResNet50_Weights.DEFAULT)

for param in aux_model.parameters():
    param.requires_grad = False

aux_model.fc = nn.Linear(2048, 12)
nn.init.kaiming_uniform_(aux_model.fc.weight, nonlinearity='relu')
aux_model.fc.requires_grad = True

```

```

optimizer = torch.optim.Adam(aux_model.parameters(), lr=0.01)
Adam_list_3 = optimizer_testing(trainloader, aux_model, optimizer, optim.
    ↪lr_scheduler.StepLR(optimizer, step_size=3), valloader)
print("Done - 3")
aux_model = resnet50(weights=ResNet50_Weights.DEFAULT)

for param in aux_model.parameters():
    param.requires_grad = False

aux_model.fc = nn.Linear(2048, 12)
nn.init.kaiming_uniform_(aux_model.fc.weight, nonlinearity='relu')
aux_model.fc.requires_grad = True
optimizer = torch.optim.Adam(aux_model.parameters(), lr=0.01)
Adam_list_4 = optimizer_testing(trainloader, aux_model, optimizer, optim.
    ↪lr_scheduler.StepLR(optimizer, step_size=4), valloader)
print("Done - 4")
aux_model = resnet50(weights=ResNet50_Weights.DEFAULT)

for param in aux_model.parameters():
    param.requires_grad = False

aux_model.fc = nn.Linear(2048, 12)
nn.init.kaiming_uniform_(aux_model.fc.weight, nonlinearity='relu')
aux_model.fc.requires_grad = True
optimizer = torch.optim.Adam(aux_model.parameters(), lr=0.01)
Adam_list_5 = optimizer_testing(trainloader, aux_model, optimizer, optim.
    ↪lr_scheduler.StepLR(optimizer, step_size=5), valloader)
print("Done - 5")

```

```

Done - 1
Done - 2
Done - 3
Done - 4
Done - 5

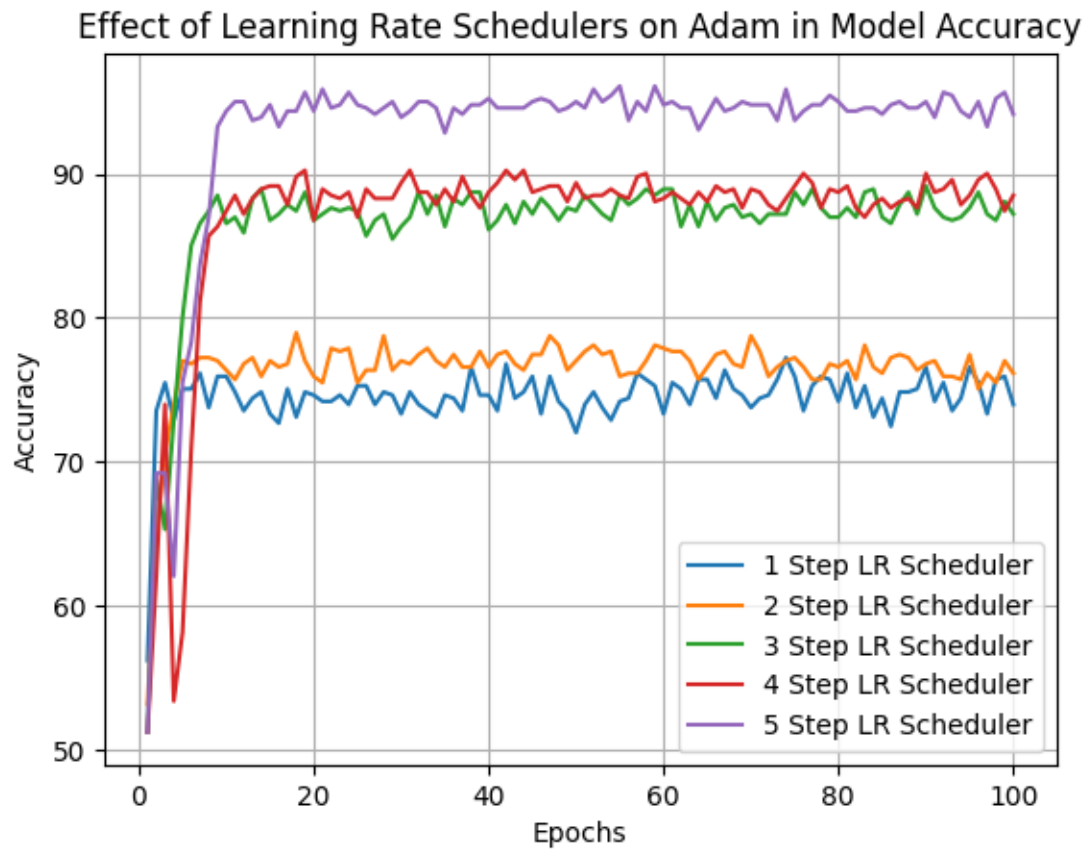
```

```

[ ]: plt.clf()
plt.plot(np.linspace(1, len(Adam_list_1), len(Adam_list_1)), Adam_list_1,
    ↪label="1 Step LR Scheduler")
plt.plot(np.linspace(1, len(Adam_list_2), len(Adam_list_2)), Adam_list_2,
    ↪label="2 Step LR Scheduler")
plt.plot(np.linspace(1, len(Adam_list_1), len(Adam_list_1)), Adam_list_3,
    ↪label="3 Step LR Scheduler")
plt.plot(np.linspace(1, len(Adam_list_1), len(Adam_list_1)), Adam_list_4,
    ↪label="4 Step LR Scheduler")
plt.plot(np.linspace(1, len(Adam_list_1), len(Adam_list_1)), Adam_list_5,
    ↪label="5 Step LR Scheduler")
plt.title("Effect of Learning Rate Schedulers on Adam in Model Accuracy")

```

```
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.grid(True)
plt.legend()
plt.show()
plt.clf()
```



<Figure size 640x480 with 0 Axes>